



TEXAS ADVANCED COMPUTING CENTER

WWW.TACC.UTEXAS.EDU



TEXAS

The University of Texas at Austin

Optimization, Profiling, & Debugging

For Modern Multicore Processors

PRESENTED BY:

Virginia Trueheart, MSIS

HPC Applications

Texas Advanced Computing Center

An Overview

Optimization & Profiling

- Write organized code
- Address architecture needs
- Address software needs
- Make good code better

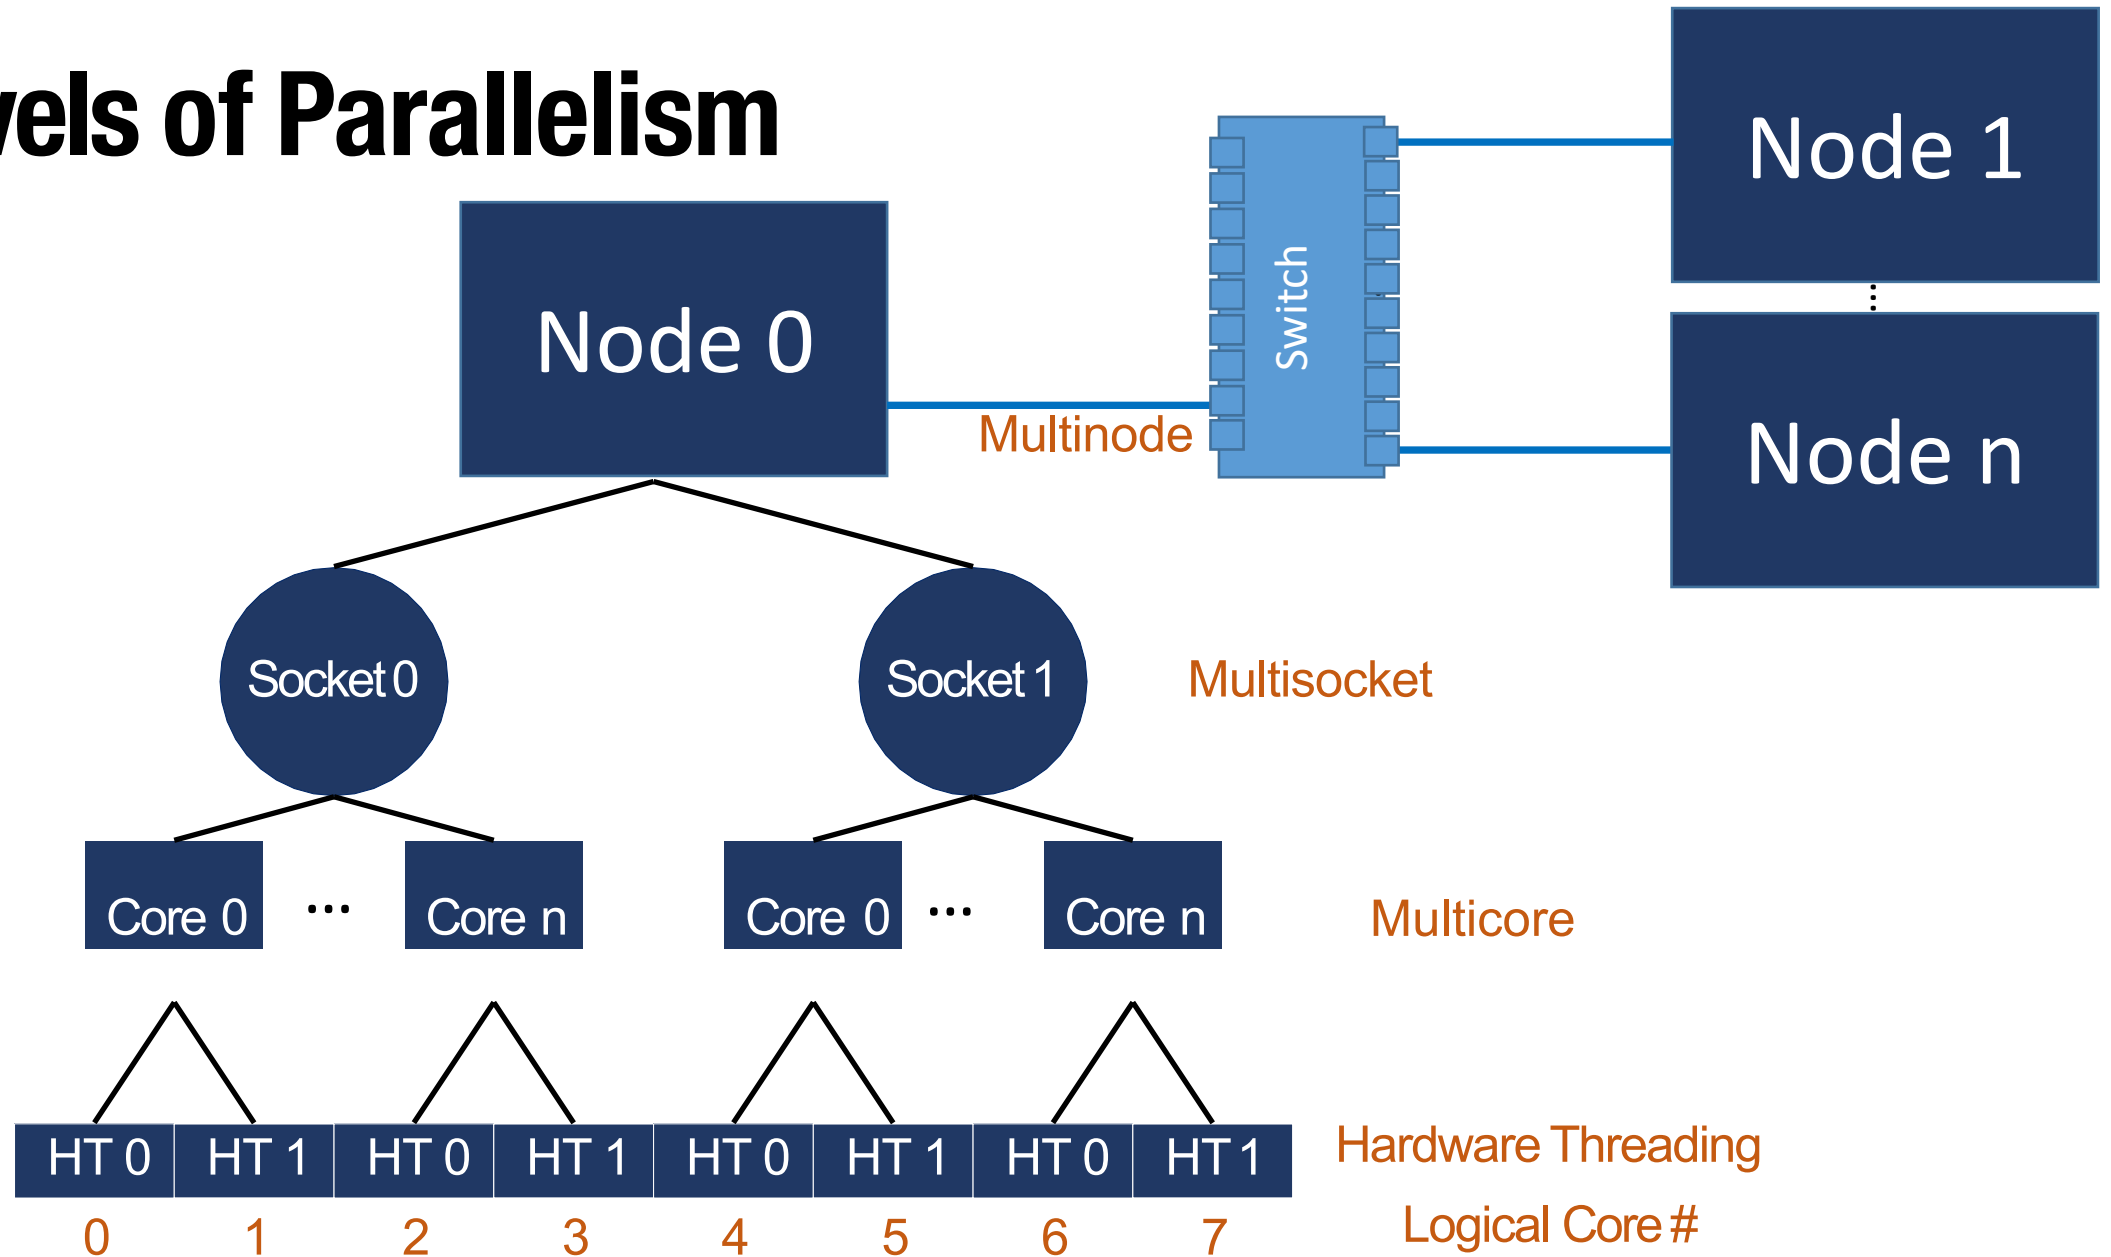
Debugging

- Find problems in your code
- Fix said problems
- Find more problems
- $(\cup \circ \square \circ) \cup \frown \text{—} \text{—} \text{—}$

Basic Optimization

- Parallelism
 - Divide computational tasks and **vectorize**
- Cache Usage
 - Smaller but **faster** storage
- Load Balancing
 - Maximize hardware and cache usage

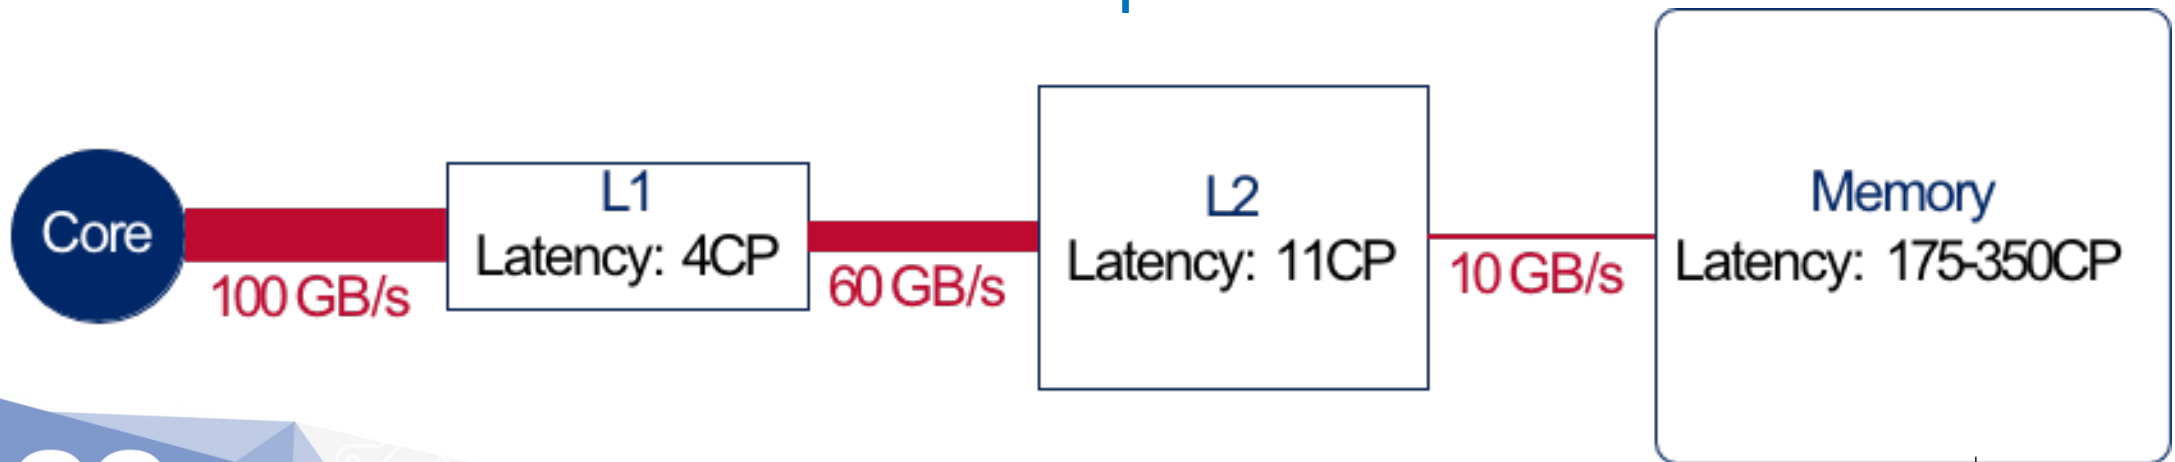
Levels of Parallelism



Caching

Memory Hierarchy


- Caches are smaller but faster than memory (higher bandwidth, less latency) than memory
- May have multiple levels of cache: L1, L2, LLC . . .
- Keeps cores fed with data
- Cache reuse is often critical to performance



Vectorization

Basics

- Cores have registers
- Data must be moved from cache/memory into registers before operating on
- To add 2 integers:
 1. Move first integer A from cache to a register
 2. Move second integer E from cache to a different register
 3. Add integers and place result in yet a different register
 4. Move result to cache/memory
- Frequencies are limited so they can only go so fast



Why Vectorization Matters

```
$ gcc vector.c -no-vec  
-o no-vec
```

```
$ ./no-vec
```

```
run time = 45.704987s
```

```
sum of a =
```

```
83886080000.000000
```

```
$ gcc vector.c -xcore-  
avx512 -qopt-zmm-  
usage=high -o avx512-  
vec-skx-zmm
```

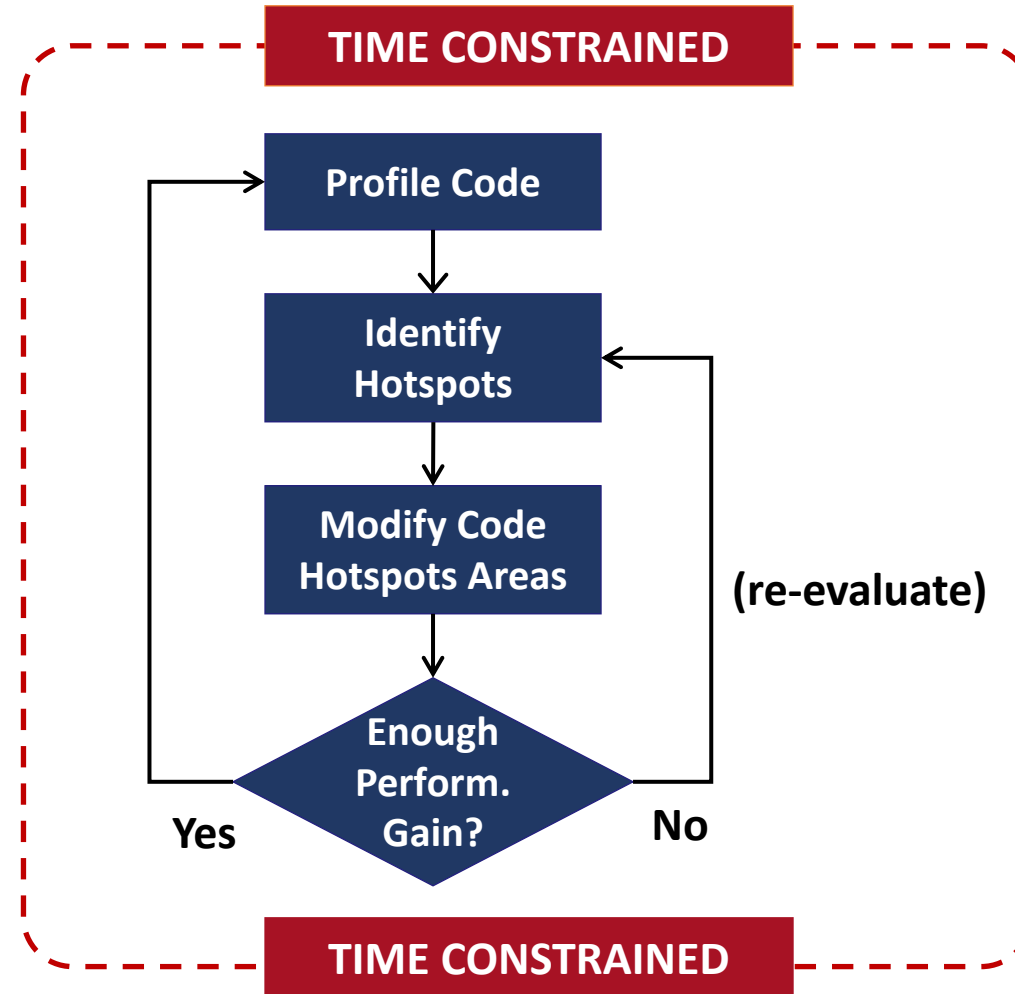
```
$ ./avx512-vec-skx-zmm
```

```
run time = 1.839410s
```

```
sum of a =
```

```
83886080000.000000
```

Optimization Improvement Process



- Iterative process
- Application dependent
- Different levels
 - Compiler Options
 - Performance Libraries
 - Code Optimizations

Profiling Basics

Controlled measurements are essential

- **Control**

Don't let the O/S decide process & memory affinity – **control it !**

- **Repeat**

No single measurement is reliable

- **Automate**

Don't try to remember how you did something – **script it !**

- **Document**

Save all these:

- ✓ Code versions
- ✓ Compilers / flags
- ✓ User inputs
- ✓ System environment
- ✓ Timers / counters
- ✓ Code output

Record Details:

Important routines should track work done and time taken

Basic Profiling Tools

CML Timer | Code Timer | gprof

Timer: Command Line

The command **time** is available in most Unix systems.

It is simple to use (no code instrumentation required).

Gives total execution time of a process and all its children in seconds.

```
$ /usr/bin/time -p ./exe
```

Elapsed wall clock time **real 9.95** ↑ **POSIX format**

Time spent in { **user mode user 9.86**
 kernel mode sys 0.06

Leave out the **-p** option to get additional information:

```
$ /usr/bin/time ./exe
```

\$ 9.860u	0.060s	0:09.95	99.9%	0+0k	0+0io	0pf+0w
User	System	Real	%CPU	Shared Unshared	Input Output	Page faults Swap

Timer: Code Section

C

```
#include <time.h>

double start, stop, time;
start = (double)clock()/CLOCKS_PER_SEC;

/* Code to time here */

stop = (double)clock()/CLOCKS_PER_SEC;
time = stop - start;
```

C (MPI)

```
#include <mpi.h>

double start, stop, time;
start = MPI_Wtime();

/* Code to time here */

stop = MPI_Wtime();
time = stop - start;
```

FORTRAN

```
INTEGER :: rate, start, stop
REAL      :: time

CALL SYSTEM_CLOCK(COUNT_RATE = rate)
CALL SYSTEM_CLOCK(COUNT = start)

! Code to time here

CALL SYSTEM_CLOCK(COUNT = stop)
time = REAL( ( stop - start ) / rate )
```

Python

```
import time

start=time.time()

# Code to time here

end=time.time()
time = end - start
```

gprof: GNU Project PROFiler

Flat Profile

- CPU time spend in each function (self and cumulative)
- Number of times a function is called
- **Useful to identify most expensive routines**

Call Graph

- Number of times a function was called by other functions
- Number of times a function called other functions
- **Useful to identify function relations**
- Suggests places where function calls could be eliminated

Annotated Source

- **Indicates number of times a line was executed**

Profiling with gprof

Use the `-pg` flag during compilation:

```
$ gcc -g -pg -o exeFile ./srcFile.c
```

```
$ gfortran -g -pg -o exeFile ./srcFile.f90
```

Run the executable. An output file `gmon.out` will be generated with the profiling information.

Execute `gprof` and redirect the output to a file:

```
$ gprof ./exeFile gmon.out > profile.txt
```

The code must
be compiled
with “-g”

```
$ gprof -l ./exeFile gmon.out > profile_line.txt Enable line-by-line profiling
```

```
$ gprof -A ./exeFile gmon.out > profile_annotated.txt Print annotated source code
```

gprof: Flat Profile

In the flat profile we can identify the most expensive parts of the code (in this case, the calls to `matSqrt`, `matCube`, and `sysCube`).

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
50.00	2.47	2.47	2	1.24	1.24	<code>matSqrt</code>
24.70	3.69	1.22	1	1.22	1.22	<code>matCube</code>
24.70	4.91	1.22	1	1.22	1.22	<code>sysCube</code>
0.61	4.94	0.03	1	0.03	4.94	<code>main</code>
0.00	4.94	0.00	2	0.00	0.00	<code>vecSqrt</code>
0.00	4.94	0.00	1	0.00	1.24	<code>sysSqrt</code>
0.00	4.94	0.00	1	0.00	0.00	<code>vecCube</code>

gprof: Call Graph Profile

index	% time	self	children	called	name
		0.00	0.00	1/1	<hicore> (8)
[1]	100.0	0.03	4.91	1	main [1]
		0.00	1.24	1/1	sysSqrt [3]
		1.24	0.00	1/2	matSqrt [2]
		1.22	0.00	1/1	sysCube [5]
		1.22	0.00	1/1	matCube [4]
		0.00	0.00	1/2	vecSqrt [6]
		0.00	0.00	1/1	vecCube [7]

		1.24	0.00	1/2	main [1]
		1.24	0.00	1/2	sysSqrt [3]
[2]	50.0	2.47	0.00	2	matSqrt [2]

		0.00	1.24	1/1	main [1]
[3]	25.0	0.00	1.24	1	sysSqrt [3]
		1.24	0.00	1/2	matSqrt [2]
		0.00	0.00	1/2	vecSqrt [6]

This table describes the call tree of the program, and was sorted by the total amount of **time** spent in each function and its children.

The lines above it list the functions that called this function, and the lines below it list the functions this one called.

gprof: -l and -A

`$ gprof -l ./exeFile gmon.out` Enable line-by-line profiling

`$ gprof -A ./exeFile gmon.out` Print annotated source code

The code must
be compiled
with “-g”

Line-by-line profiling (-l)

Flat profile:

Each sample counts as 0.01 seconds.

%	cumulative	self		self	total	
time	seconds	seconds	calls	Ts/call	Ts/call	name
100.00	47.21	47.21				StaticFunc
(gprof_test.c:23 @ 40195a)						
0.00	47.21	0.00	100	0.00	0.00	StaticFunc
(gprof_test.c:18 @ 401927)						
0.00	47.21	0.00	1	0.00	0.00	TestFunc
(gprof_test.c:7 @ 4018c8)						

Line
number

```
#include<stdio.h>
void TestFunc();
static void StaticFunc();
void TestFunc()
1-> {
    int i = 0;
    printf("In TestFunc\n");
    for (i=0; i<100; i++)
        StaticFunc();
}
static void StaticFunc()
100-> {
    int i = 0;
    printf("In StaticFunc\n");
    for (i=0; i<100000000; i++);
}
int main(void)
##### -> {
    printf("In main\n");
    TestFunc();
    return 0;
}
```

Top 10 Lines:

Line	Count
18	100
7	1

Execution Summary:

3	Executable lines in this file
3	Lines executed
100.00	Percent of the file executed
101	Total number of line executions
33.67	Average executions per line

Annotated source listing (-A)

Advanced Profilers

When gprof isn't enough you can consider some of the following tools:

- IPM (Integrated Performance Monitoring)
- HPCToolKit
- TAU
- VTune

Tutorials are available online for many of these and they can produce far more information about the jobs and help you generate graphs and other visualizations.

- Free
- Paid (but the viewer can be downloaded for free to your desktop)

Basic Debugging

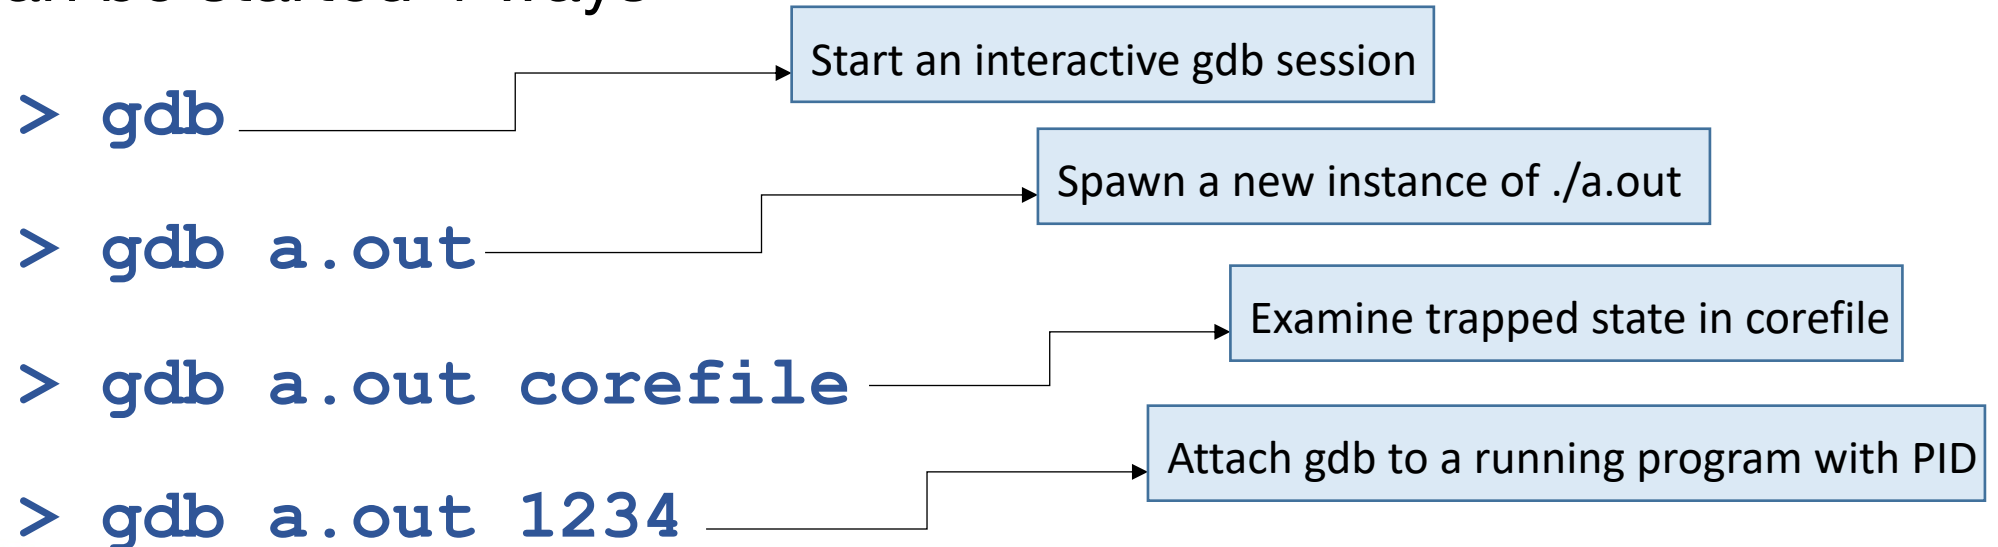
- Show **backtraces**
View the history of the code up to the point of interest
- Set **breakpoints**
Pause your code in various states
- Display the value of individual **variables**
Like print statements, this allows you to see the state of the code
- Set new values
 - Controlled values make mistakes easy to see
- Run individual **steps** of the program
 - Ensure that every function and subsection runs properly independently

gdb: GNU Debugger

As with the profile, recompile with the **-g** flag

```
> gcc -g -o hello hello.c
```

gdb can be started 4 ways



hello.c

```
#include <stdio.h>

void foo();

int main ()
{
    printf("inside main\n");
    foo();
    return;
}

void foo ()
{
    int i, total=0;
    printf("inside foo\n");
    for (i=0; i<1000; i++)
        total += i;
}
```

gdb Commands

run – starts program

print – print a variable located in current scope

next – executes current command and moves to next command

break - set a break point

continue – run until next break point or termination

delete – delete a break point

condition – make a break point conditional

where – show current function in stack trace

An Example

```
$ gdb ./hello
```

```
GNU gdb (GDB) Red Hat Enterprise Linux  
...<http://www.gnu.org/software/gdb/bugs/>
```

```
Reading symbols from  
/scratch/03658/vtrue/training/petasc19/hello...done
```

```
(gdb) run
```

```
Starting program:  
/scratch/03658/vtrue/training/petasc19/./hello
```

```
inside main
```

```
inside foo
```

Cont.

```
(gdb) break main
```

```
Breakpoint 1 at 0x4005ae: file hello.c, line 6.
```

```
(gdb) run
```

```
Starting program:
```

```
/scratch/03658/vtrue/training/petascade19/./hello
```

```
Breakpoint 1, main () at hello.c:6
```

```
6    printf("inside main\n");
```


Other Kinds of Bugs

Floating Point Errors

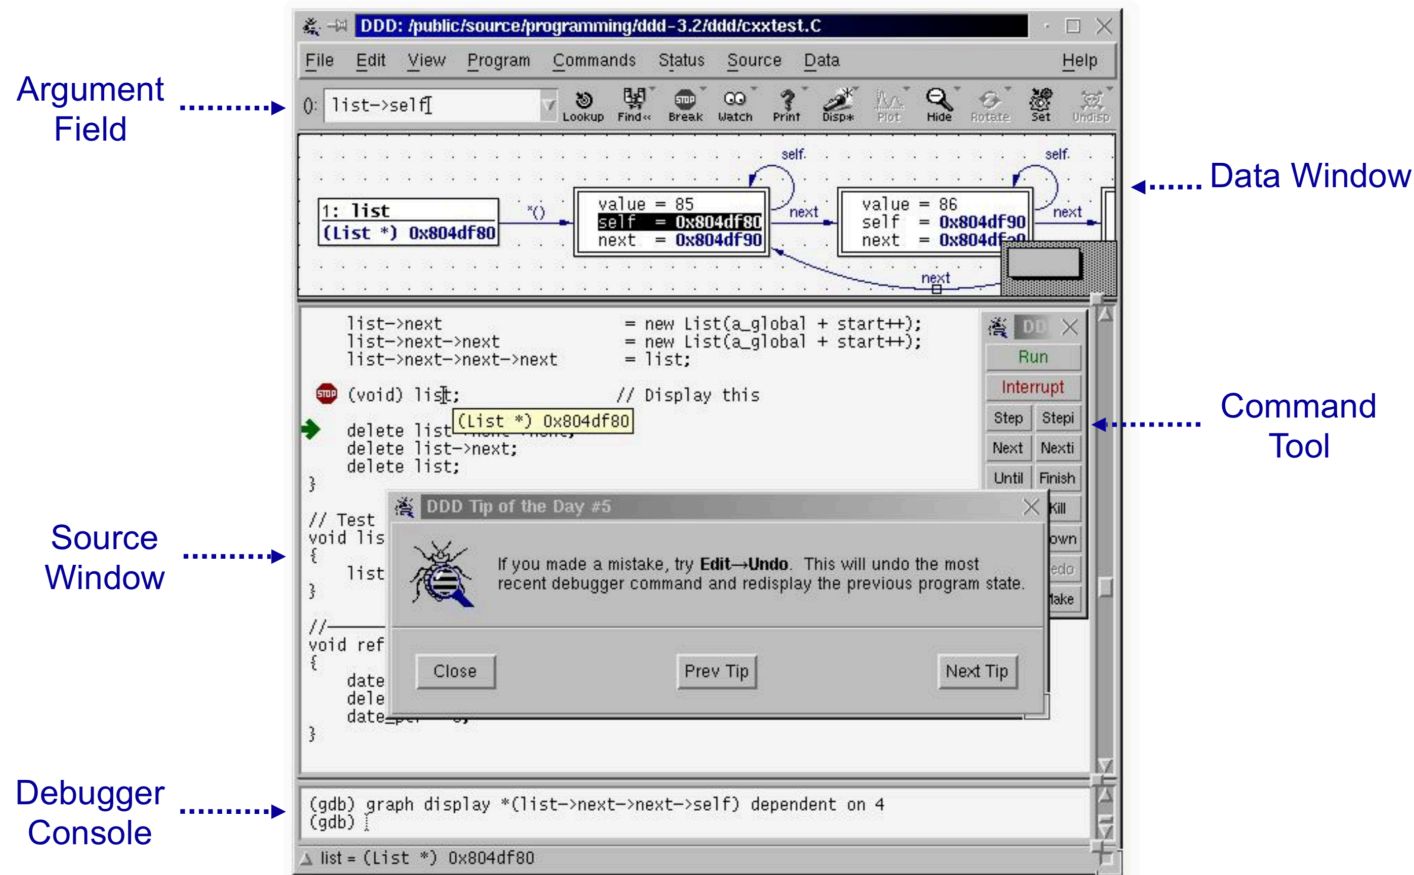
- Result in **NaN** (not-a-number) issues
- Different compilers may have different outputs
- Test for these periodically to preemptively trap a diverged solution

Memory Errors

- Difficult to trap: use **glibc**, **dmalloc**, **Electric Fence**, or **Valgrind** to help
- These runtime checks will slow down the performance of the code
- Restrict memory checks to non-production runs

Advanced Debuggers

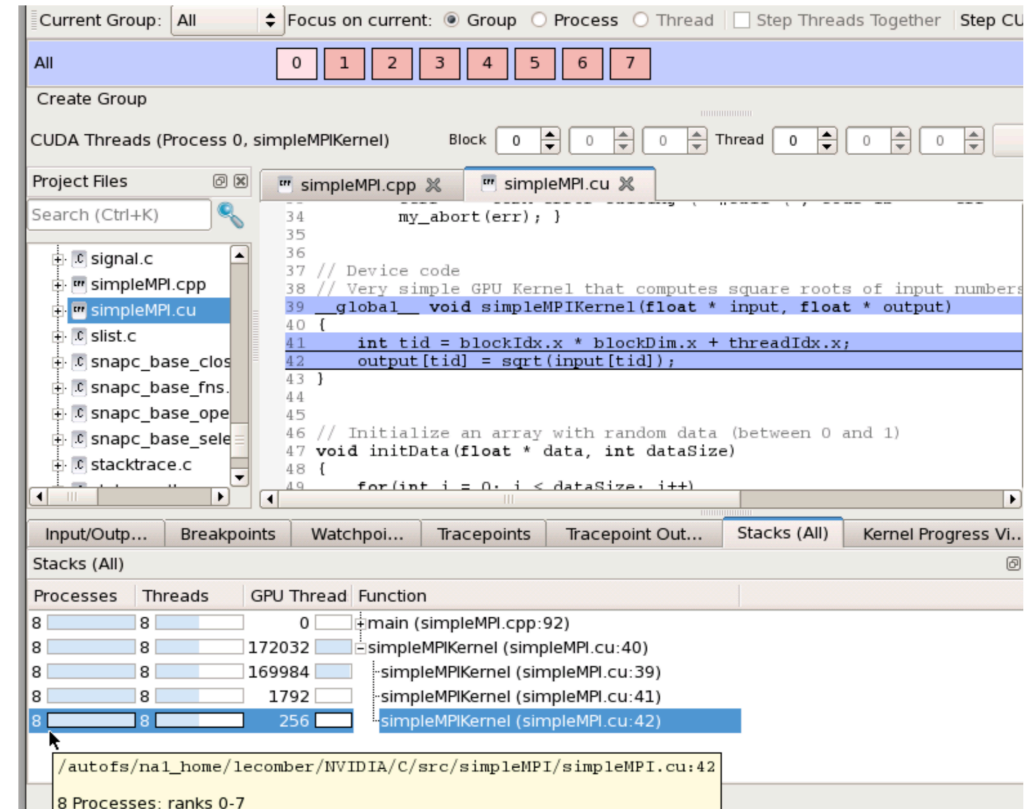
Data Display Debugger – DDD



Cont.

Arm DDT

- This is a commercial debugger but is available on some systems like Stampede2
- Can interrogate over 100k MPI tasks
- Includes some basic memory debugging tools



Resources

Vectorization -

<https://learn.tacc.utexas.edu/mod/page/view.php?id=48>

gprof - https://web.eecs.umich.edu/~sugih/pointers/gprof_quick.html
& <https://sourceware.org/binutils/docs/gprof/>

gdb - <http://www.sourceware.org/gdb/current/onlinedocs/gdb.html>

DDD - <http://www.gnu.org/manual/ddd/>

License

©The University of Texas at Austin, 2019

This work is licensed under the Creative Commons Attribution Non-Commercial 3.0 Unported License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/>

When attributing this work, please use the following text: “Introduction to Many Core Programming”, Texas Advanced Computing Center, 2018. Available under a Creative Commons Attribution Non-Commercial 3.0 Unported License.

